

Homework # 3

Mitchell Scott
(mtscot4)

1. **Energy appliance regression** (50 pts, written + code):

We again use the Energy dataset (`energydata.zip`) from Homework #1. As a reminder, each sample contains measurements of temperature and humidity sensors from a wireless network, weather from a nearby airport station, and the recorded energy use of lighting fixtures to predict the energy consumption of appliances in a low energy house. Put all code into `q1.py`. There are 2 classes: **FeatureSelection** and **Regression**. Complete implementation of the following class methods. You may use helper methods from the `sklearn.feature_selection` module and regressors from `sklearn.linear_model`.

- (a) **Code:** `FeatureSelection.rank_correlation(x, y)` that takes in a numpy 2d array, \mathbf{x} , a numpy 1d array, \mathbf{y} , and returns the rank of the features in \mathbf{x} based on Pearson correlation of each individual feature to the target value. The function should return a numpy array in descending order of the most correlated feature column (i.e., `[3, 1, 0, 2]` would mean that the 4th feature has the highest correlation to \mathbf{y} followed by 2nd, then 1st, and lastly 3rd feature.)

Solution. Implimented in `q1.py`.

- (b) **Code:** `FeatureSelection.Lasso(x, y)` that takes in a numpy 2d array, \mathbf{x} , a numpy 1d array, \mathbf{y} , and returns the rank of the features in \mathbf{x} based on coefficients of Lasso regression (ℓ^1 regularization). The function should return a numpy array in descending order of the absolute magnitude of Lasso regression coefficient. do not include features with zero Lasso coefficient. For example, for 4-features input \mathbf{X} , if Lasso coefficients are `[2.3, -4.5, 0, 7.6]`, we should return `[3, 1, 0]`, to mean that the 4th feature has the highest coefficient followed by 2nd, then 1st, and the 3rd feature is dropped. Note that the returned array is likely to be shorter than the original dimension of \mathbf{X} .

Solution. Implimented in `q1.py`.

- (c) **Code:** `FeatureSelection.stepwise(x, y)` that takes in a numpy 2d array, \mathbf{x} , a numpy 1d array, \mathbf{y} , and returns the rank of the features in \mathbf{x} based by greedy adding one remaining feature at a time to the set. To do this, we try adding one remaining feature at a time to the set, train a linear regressor, and select next feature to add based on the highest decrease in RMSE. Stop the process when RMSE does not decrease. For simplicity, we do this process on the training set \mathbf{X} , but feel free to use cross validation for this step. See slides for details or read about reference implementation `SequentialFeatureSelector`. Return ranking of features in descending order of importance. Note that the returned array is likely to be shorter than the original dimension of \mathbf{X} .

Solution. Implimented in `q1.py`.

- (d) **Written:** report up to 10 first features selected by each method 1a,1b, and 1c in a table. Comment on overlap and differences in the resulting feature sets.

Solution. For the `rank_correlation` method, we arrived at [24, 23, 8, 10, 9, 20, 6, 17, 22, 2]. For the `lasso` method, we arrived at, [12, 22, 21, 20, 19, 17, 15, 14, 13, 24]. Lastly, For the `stepwise` method, we arrived at, [12, 22, 20, 19, 17, 15, 14, 10, 24, 7]. We see that `lasso` and `stepwise` give rather similar results with the first 7 out of 8 entries being the same. Additionally, we see 24 represented close to the bottom in both methods as well. Comparing either `lasso` or `stepwise` with `rank_correlation`, we see that the results are not all that similar. Of course all methods have 24, 20, 17, and 22, but whereas, `lasso` starts with 22,20, and 17, `rank_correlation` has those entries at the end. Conversely, `stepwise` has 24 as one of the last, but `rank_correlation` has it first.

Additionally, while `lasso` and `stepwise` have very similar results, `stepwise` took noticeably longer to compute.

- (e) **Code:** `Regression.Ridge(train_x, train_y, test_x, test_y)`: Implement Ridge regression and returns prediction on test set similar to HW1. You may use `sklearn` for the regressor implementation

Solution.

- (f) **Written:** use your code in 1e to train and test Ridge regression on the whole feature set (“No selection”), and on top 10 features selected in 1a, 1b, 1c. Report in a table the RMSE and R^2 of the Ridge regressor on the validation dataset and test dataset. Comment on differences/similarities of performance of Ridge regression under the different feature selection methods.

Solution.

- (g) **Code:** `Regression.DecisionTreeRegressor (train_x, train_y, test_x, test_y, max_depth, min_items)`: Implement Regression Tree training and prediction, and return prediction on validation and test sets similar to HW1. You may use `sklearn` for the underlying regressor tree implementation.

Solution.

- (h) **Written:** tune the hyperparameters of your regression tree in 1g on whole feature set by varying the `max_depth` and `min_items` and selecting the best configuration based on the validation set. Report the best configuration and RMSE and R^2 obtained on validation set.

Solution. max depth = 2 min=1 'train-rmse': 0.9710814885538052, 'train-r2': 0.11075240052065438, 'test-rmse': 0.9263770467806784, 'test-r2': -0.05816567665070438
max = 2, min = 3 'train-rmse': 0.9710814885538052, 'train-r2': 0.11075240052065438, 'test-rmse': 0.9263770467806786, 'test-r2': -0.0581656766507046
max = 2, min = 5 'train-rmse': 0.9710814885538052, 'train-r2': 0.11075240052065438, 'test-rmse': 0.9263770467806784, 'test-r2': -0.05816567665070438
max depth = 5 min 1 Tree Regressor: 'train-rmse': 0.9262768557261435, 'train-r2': 0.26385519598795015, 'test-rmse': 1.0116562187956508, 'test-r2': -0.5049925470028667
max depth =5, min = 3 'train-rmse': 0.9253569559487326, 'train-r2': 0.2667751494724081, 'test-rmse': 1.0116562187956508, 'test-r2': -0.5049925470028664
max = 5, max =5 'train-rmse': 0.9261954052795444, 'train-r2': 0.2641140879791416, 'test-rmse': 1.0116562187956508, 'test-r2': -0.5049925470028664
max = 8, min

```
=1 'train-rmse': 0.8304318880014541, 'train-r2': 0.5244282269620424, 'test-
rmse': 1.07060762395057, 'test-r2': -0.8876603515468597 max = 8, min 3
'train-rmse': 0.8316872965740683, 'train-r2': 0.5215459093942053, 'test-rmse':
1.0294738928507934, 'test-r2': -0.6138526404431388 max = 8, min 5 'train-
rmse': 0.8413936190345813, 'train-r2': 0.498816394208888, 'test-rmse': 1.0220726960471527,
'test-r2': -0.5679408438248181
```

- (i) **Written:** Train the regressor using best configuration on full dataset as training (train=train+validation), and report RMSE and R^2 on test dataset. Compare performance of your tree regressor to the best performance of Ridge regression in 1f.

Solution.

- (j) **Written:** visualize top 3 levels in your decision tree in 1h and 1i. Compare the top-level features used by your tree, with the features selected in 1d. Comment on overlap/order of importance of features in the two methods.

2. **Spam classification using Naïve Bayes, Random Forest, and GBDT** (50 pts, written + code): We will use the same email spam dataset from HW2 the email spam dataset, which contains 4601 e-mail messages that have been split into 3000 training (spam.train.dat) and 1601 test emails (spam.test.dat). 57 features have been already extracted for you, with a binary label in the last column. All the specified functions should be in the file 'q2.py'.

- (a) **Code:** Write a Python function `eval_randomforest(trainx, trainy, testx, testy, num_trees, max_depth, min_items)` that fits a Random Forest model to the training data. You can use sklearn module for this part. The function should accept as input numpy 2d arrays, and return a dictionary containing the accuracy and AUC for the training and test sets and the predicted probabilities for the test set: `return {"train-acc": train_acc, "train-auc": train_auc, "test-acc": test_acc, "test-auc": test_auc, "test-prob": test_prob}`. The values for accuracy and AUC should be scalar numeric values, while test-prob should be a numpy 1-d array with the predicted probability for the positive class 1, for the test. Same format as HW2.

Solution.

- (b) **Written:** Tune RandomForest for the Spam classification problem by optimizing the hyperparameters `num_trees`, `max_depth`, `min_items` based on cross validation over the training set only. Report best configuration, and average AUC and Accuracy across validation folds.

Solution.

- (c) **Written:** Train Random forest on all training data with using best hyperparameters identified in 1c, and report AUC, F1, and Accuracy on the test set. Compare in a table to the performance to NB and LR from Homework 2.

Solution.

- (d) **Written:** Report top 10 most important features for NB and RF (describe your criteria chosen for "most important" in each case).

Solution.

- (e) **Code:** Write a Python function `eval_gbdtd(trainx, trainy, testx, testy, num_estimators, learning_rate)` that fits a `GradientBoostingClassifier` model to the training data. You can use `sklearn` module for this part. The function should accept as input numpy 2d arrays, and return a dictionary containing the accuracy and AUC for the training and test sets and the predicted probabilities for the test set: `return {"train-acc": train_acc, "train-auc": train_auc, "test-acc": test_acc, "test-auc": test_auc, "test-prob": test_prob}`. The values for accuracy and AUC should be scalar numeric values, while test-prob should be a numpy 1-d array with the predicted probability for the positive class 1, for the test. Same format as HW2.

Solution.

- (f) **Written:** Tune GBDT for the Spam classification problem by optimizing the hyperparameters `num_estimators` and `learning_rate` at least (you may choose to tune others). Use cross validation over the training set only. Report the best configuration, and the average AUC, F1, and Accuracy across validation folds.

Solution.

- (g) **Written:** Train GBDT on all training data with your best hyperparameters, and report AUC, F1, and Accuracy on the test set. Compare in a table to the performance of RF and NB. Comment on relative improvements / performance for this problem by the three different models.

Solution.

Acknowledgements

I would like to acknowledge that I attended both TA Swati's and TA Ziyang's office hours over the past week.